

hi:

```
hi: 0.980584  
sup: 0.019416
```

sup:

```
hi: 0.150009  
sup: 0.849991
```

The notebook had one training setup issue, it treated a two-class one-hot classification problem like a regression problem by using `loss='mse'` and `metrics=['mae']`.

I corrected the compile step to use `categorical_crossentropy` with `accuracy`, which matches the softmax output layer and the one-hot gesture labels used in the dataset.

After changing the training objective, I updated the downstream plotting code so it uses `accuracy` and `val_accuracy` instead of `mae` and `val_mae`.

I also changed the plot title, y-axis label, legend text, and history-key references so the notebook's visuals match the classification task instead of a regression-style metric.